



Sistemas de Operação / Fundamentos de Sistemas Operativos

File systems in a nutshell

Artur Pereira <artur@ua.pt>

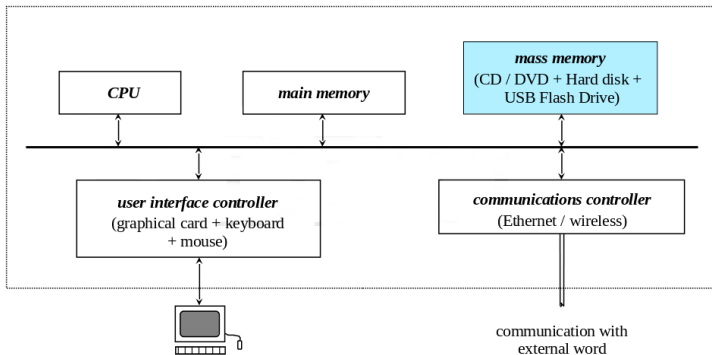
DETI / Universidade de Aveiro

Outline

- 1 Overview
- 2 Mass storage
- 3 The file concept
- 4 Inodes

Overview

- Simple view of a computational system, highlighting the mass memory component:
- **File system** is the part of the operating system responsible to manage access to mass storage



Mass storage


Types of mass storage devices

Type	Technology	Capacity (Gbytes)	Type of use	Transfer rate (Mbytes/s)
CD-ROM	mechanical / optical	0.7	read	0.5
DVD	mechanical / optical	4–8	read	0.7
HDD	mechanical / magnetical	250–4000	read / write	480
USB FLASH	semiconductor	2–256	read / write	60(r) / 30(w)
SSD	semiconductor	64–512	read / write	500

Mass storage

Operational abstraction of mass storage

- **Mass storage** can be seen in operational terms as a very simple model
 - each device is represented by an array of NTBK storage blocks, each one consisting of BKSZ bytes (typically BKSZ ranges between 256 and 32K)
 - access to each block for reading or writing can be done in a random manner
- This is called **Logical Block Addressing – LBA**
 - Blocks are located by an integer index (0, 1, ...)
 - The ATA Standard included 22-bit LBA, 28-bit LBA, and 48-bit LBA




- **Note that:**
 - a block is the only unit of interaction
 - thus, a single byte **can not** be accessed directly
-
- What to do to change a byte of a block?

Mass storage

User abstraction of mass memory

- Some considerations:
 - Despite creating a uniform model, LBA is **not an appropriate way** for a user to access mass memory data
 - Direct manipulation of the information contained in the physical device **can not be left** entirely to the responsibility of the application programmer
 - Access must be guided by quality criteria, in terms of efficacy, efficiency, integrity and sharing
- Thus, a **uniform model of interaction** is required



- **Solution:** the **file concept**

File concept

What is a file?

- **file** is the logical unit of storage in mass memory
 - meaning that reading and writing information is always done within the strict scope of a file
 - Remember that physically the unit of interaction is the block
- Basic elements of a file:
 - **identity name/path** – the (generic) way of referring to the information
 - **identity card** – meta-data (owner, size, permissions, times, ...)
 - **contents** – the information itself, organized as a sequence of bits, bytes, lines or registers, whose precise format is defined by the creator of the file and which has to be known by whoever accesses it
- From the point of view of the application programmer, a file is understood as an abstract data type, characterized by a set of **attributes** and a set of **operations**

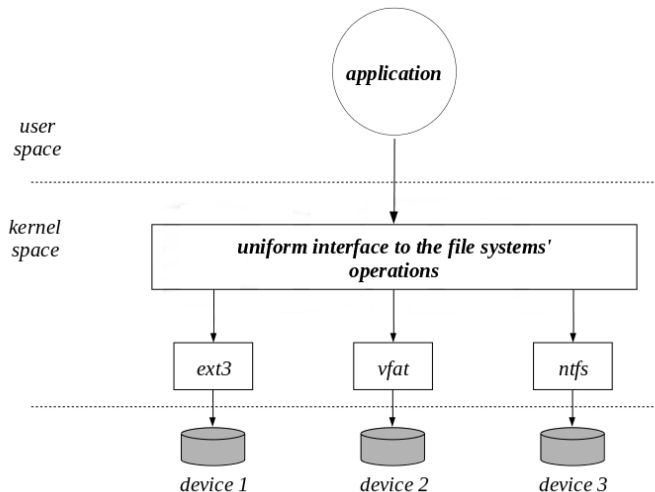
File concept

Role of operating system

- A role of the operating system is to implement this data type, providing a set of operations (**system calls**) which establishes a simple and secure communication interface for accessing the mass memory
- The **file system** is part of the operating system dedicated to this task
- Different implementations of the file data type lead to different types of file systems
 - Ex: ext3, FAT, NTFS, APFS, ISO 9660, ...
- Nowadays, a single operating system implements different types of file systems, associated with different physical devices, or even with the same
 - This feature facilitates interoperability, establishing a common means of information sharing among heterogeneous computational systems

File concept

Virtual file system



File concept

Types of files

- From the operating system point of view, there are different types of files:
 - **ordinary/regular file** – file whose contents is of the user responsibility
 - **directory** – file used to track, organize and locate other files and directories
 - **shortcut (symbolic link)** – file that contains a reference to another file (of any type) in the form of an absolute or relative path
 - **character device** – file representing a device handled in bytes
 - **block device** – file representing a device handled in blocks
 - **socket** – file used for inter-process and inter-machine communication
 - **named pipe** – file used for inter-process communication
- Note that text files, image files, video files, application files, etc., are all **regular files**

File concept

Attributes of files


- Common attributes of a file
 - **type** – one of the referred above
 - **name/path** – the way users usually refer to the file
 - **internal identification** – the way the file is known internally
 - **size(s)** – size in bytes of information; space occupied on disk
 - **ownership** – who the file belongs to
 - **permissions** – who can access the file and how
 - **access and modification times** – when the file was last accessed or modified
 - **location of data in disk** – ordered set of blocks/clusters of the disk where the file contents is stored

-
- Remember that a disk is a set of numbered blocks

File concept

Operations on files (1)

- Common operations on regular files
 - **creation**, **deletion**
 - **opening**, **closing** – direct access is not allowed
 - **reading**, **writing**, **resizing**
 - **positioning** – in order to allow random access



File concept

Operations on files (2)

- Common operations on directories
 - **creation**, **deletion** (if empty)
 - **opening** (only for reading), **closing**
 - **reading** (directory entries)
 - A directory can be seen as a set/sequence of (directory) entries, each one representing a file (of any valid type)
- Common operations on shortcuts (symbolic links)
 - **creation**, **deletion**
 - **reading** (the value of the symbolic link)
- Common operations on files of any type
 - **get attributes** (access and modification times, ownership, permissions)
 - **change attributes** (access and modification times, permissions)
 - **change ownership** (only root or admin)

File concept

Typical file operations on Unix

- As referred to before, the operations are based on **system calls**
- system calls common to any type of file
 - **close**, **mknod**, **chmod**, **chown**, **stat**, **utimes**, ...
- system calls on regular files
 - **creat**, **open**, **link**, **unlink**, **read**, **write**, **truncate**, **lseek**, ...
- system calls on directories
 - **mkdir**, **rmdir**, **getdents**, ...
- system calls on symbolic links
 - **readlink**, **symlink**, ...

-
- On a terminal execute `man 2 <<syscall>>` to see a description

Inodes

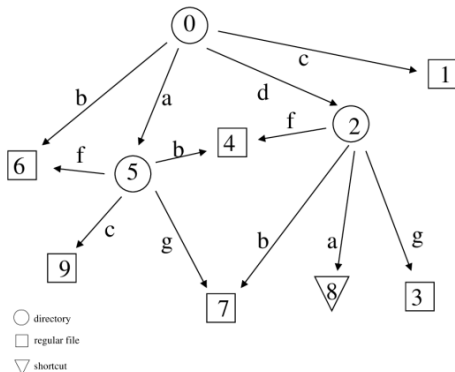
What is an inode?

- In Unix, the **inode** (identification node) plays a central role in the implementation of the file data type
 - An inode is typically identified by an integer number
 - It corresponds to the **identity card** of a file and contains:
 - file type
 - owner information
 - file access permissions
 - access times
 - file size (in bytes and blocks)
 - sequence of disk blocks with the file contents
 - The **name/path** is not in the inode
 - it is in the directory entry
-
- disk inodes vs. in-core inodes

Inodes

Hierarchy of files

- Every file uses one and only one inode
- Same inode can have different pathnames
- Hierarchy of files **may not be a tree**



- The contents of a disk can be seen as a graph, where
 - Nodes are the files (directories, regular files, shortcuts, ...), each one having an associated inode
 - Arrows define the hierarchy
- What is a directory?
- What is a link?
- What is a shortcut (symlink)?